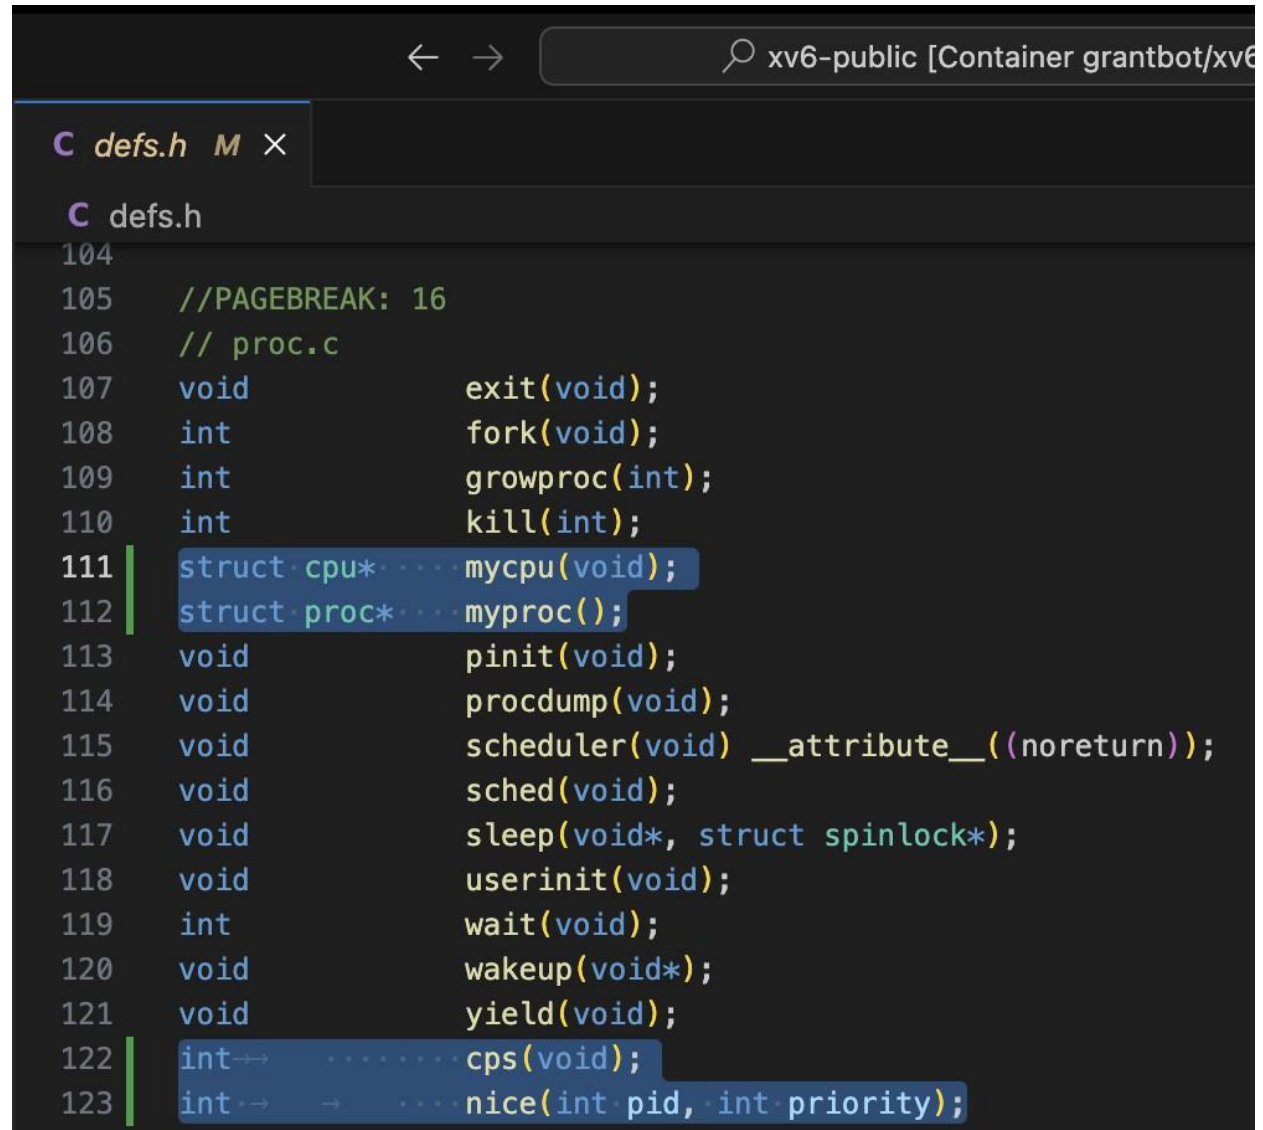


Explanation

Task 1 – Nice System Call Implementation

Changes in files:

Defs.h



```
104
105 //PAGEBREAK: 16
106 // proc.c
107 void      exit(void);
108 int       fork(void);
109 int       growproc(int);
110 int       kill(int);
111 struct cpu* mycpu(void);
112 struct proc* myproc();
113 void      pinit(void);
114 void      procdump(void);
115 void      scheduler(void) __attribute__((noreturn));
116 void      sched(void);
117 void      sleep(void*, struct spinlock*);
118 void      userinit(void);
119 int       wait(void);
120 void      wakeup(void*);
121 void      yield(void);
122 int       cps(void);
123 int       nice(int pid, int priority);
```

```
int .....lapicid(void);
```

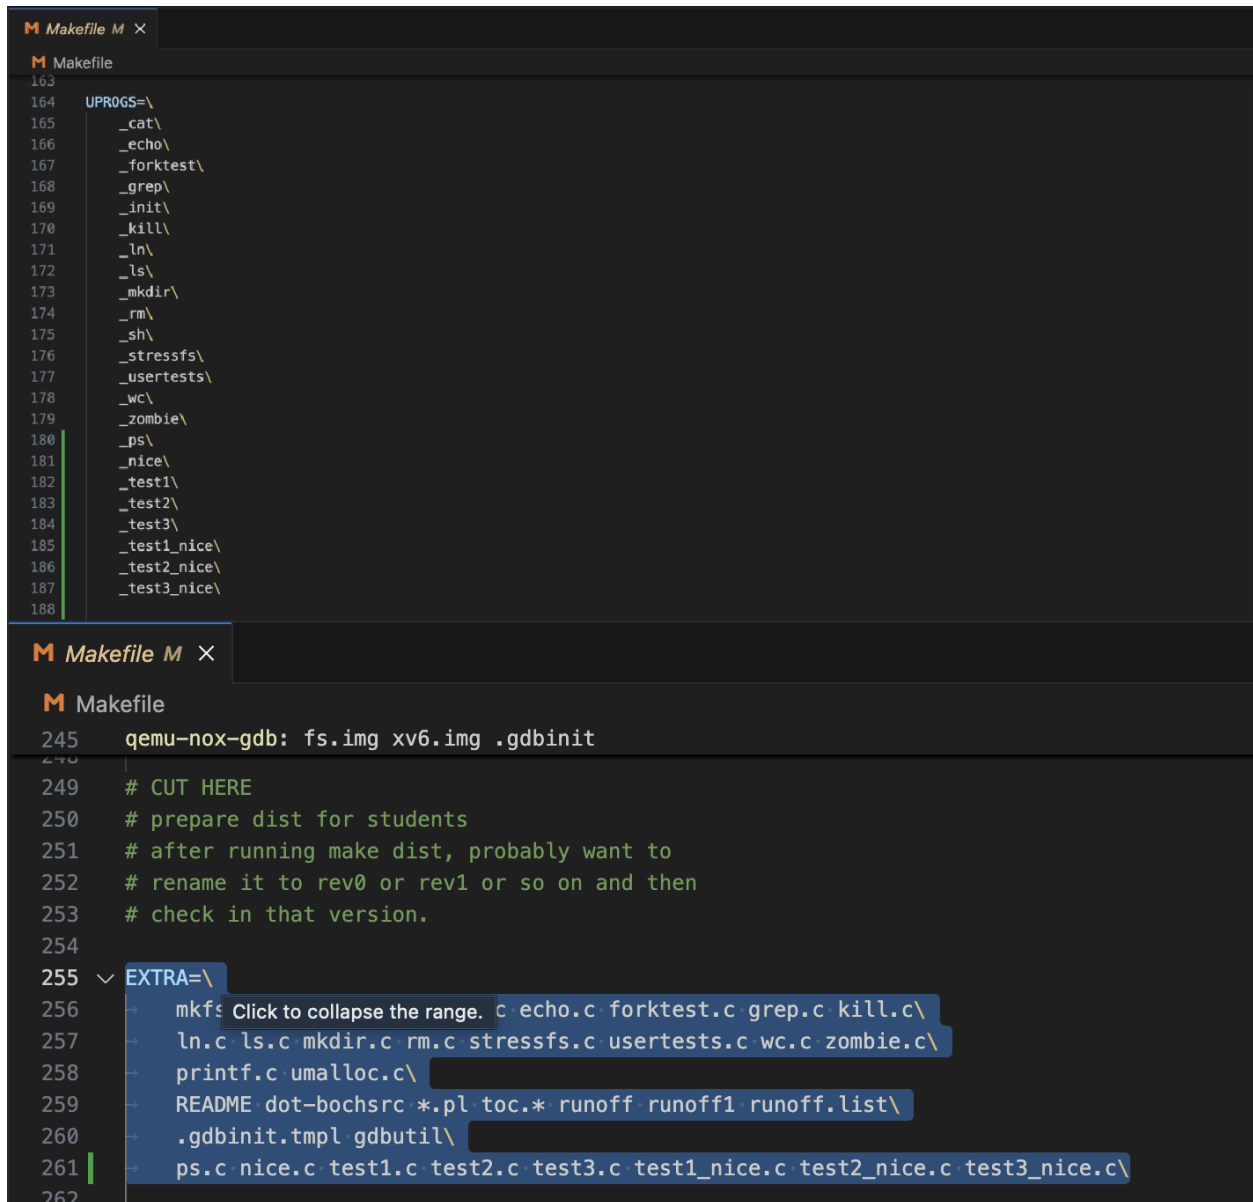
Added function definitions in defs.h

lapic.c

```
C lapic.c M X
C lapic.c
39  #define ERROR    (0x0370/4)    // Local Vector Table 3 (ERROR)
40  #define MASKED    0x00010000    // Interrupt masked
41  #define TICR      (0x0380/4)    // Timer Initial Count
42  #define TCCR      (0x0390/4)    // Timer Current Count
43  #define TDCR      (0x03E0/4)    // Timer Divide Configuration
44
45  volatile uint *lapic; // Initialized in mp.c
46
47  static void
48  lapicw(int index, int value)
49  {
50      lapic[index] = value;
51      lapic[ID]; // wait for write to finish, by reading
52  }
53  //PAGEBREAK!
54
55  int
56  lapicid(void)
57  {
58      if (!lapic)
59          return 0;
60      return lapic[ID] >> 24;
61  }
62
63  void
```

The above function is used in mycpu() function.

MAKEFILE



The image shows two screenshots of a Makefile editor. The top screenshot displays the `UPROGS` list, which includes various system utilities and test programs. The bottom screenshot displays the `EXTRAS` list, which includes source files for the same utilities and test programs. The `EXTRAS` list is expanded, showing a list of files with a tooltip that says "Click to collapse the range."

```
163
164 UPROGS=\
165     _cat\
166     _echo\
167     _forktest\
168     _grep\
169     _init\
170     _kill\
171     _ln\
172     _ls\
173     _mkdir\
174     _rm\
175     _sh\
176     _stressfs\
177     _usertests\
178     _wc\
179     _zombie\
180     _ps\
181     _nice\
182     _test1\
183     _test2\
184     _test3\
185     _test1_nice\
186     _test2_nice\
187     _test3_nice\
188

245 qemu-nox-gdb: fs.img xv6.img .gdbinit
246
249 # CUT HERE
250 # prepare dist for students
251 # after running make dist, probably want to
252 # rename it to rev0 or rev1 or so on and then
253 # check in that version.
254
255 EXTRA=\
256     mkfs.c echo.c forktest.c grep.c kill.c\
257     ln.c ls.c mkdir.c rm.c stressfs.c usertests.c wc.c zombie.c\
258     printf.c umalloc.c\
259     README dot-bochsrc *.pl toc.* runoff runoff1 runoff.list\
260     .gdbinit.tmpl gdbutil\
261     ps.c nice.c test1.c test2.c test3.c test1_nice.c test2_nice.c test3_nice.c\
262
```

Added files under UPROGS and EXTRAS to get them into simulation env.

nice.c

```
C nice.c U X
C nice.c
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4  #include "fcntl.h"
5  #include "param.h"
6
7  int
8  main(int argc, char *argv[])
9  {
10     int priority, pid=-1;
11     // old_priority=-1;
12     if(argc < 2 || argc > 3){
13         printf(2,"Usage: nice pid priority\n");
14         exit();
15     }
16     if(argc==2){
17         //change the priority of the current process to the priority given in prompt
18         priority = atoi(argv[1]);
19         if (priority < 0 || priority > 5){
20             printf(2,"Invalid priority (0-5)!\n");
21             exit();
22         }
23         int current_pid=getpid();
24         int res=nice(current_pid,priority);
25         printf(1,"%d %d\n",res/MOD,res%MOD);
26         exit();
27     }
28     else{
29         //if three arguments are given
30         priority = atoi(argv[2]);
31         if (priority < 0 || priority > 5){
32             printf(2,"Invalid priority (0-5)!\n");
33             exit();
34         }
35         pid = atoi(argv[1]);
36         int res=nice(pid,priority);
37         printf(1,"%d %d\n",res/MOD,res%MOD);
38         exit();
39     }
40 }
```

The program modifies the priority of a process in a Unix-like system. If only a priority is given, it changes the current process's priority. If a process ID (pid) and priority are provided, it changes the priority of that specified process. It validates inputs to ensure the priority is within the acceptable range (0-5) and calls a `nice()` function to perform the change.

params.h

```
C param.h M X
C param.h
1  #define NPROC      64 // maximum number of processes
2  #define KSTACKSIZE 4096 // size of per-process kernel stack
3  #define NCPU       8 // maximum number of CPUs
4  #define NOFILE     16 // open files per process
5  #define NFILE      100 // open files per system
6  #define NINODE      50 // maximum number of active i-nodes
7  #define NDEV        10 // maximum major device number
8  #define ROOTDEV     1 // device number of file system root disk
9  #define MAXARG      32 // max exec arguments
10 #define MAXOPBLOCKS 10 // max # of blocks any FS op writes
11 #define LOGSIZE      (MAXOPBLOCKS*3) // max data blocks in on-disk log
12 #define NBUF         (MAXOPBLOCKS*3) // size of disk block cache
13 #define FSSIZE       1000 // size of file system in blocks
14 #define PLEVELS      5 // Levels of priority
15 #define MOD           65 // to extract pid and priority
16
```

PLEVELS is the priority levels and MOD is used to get PID and old value using dividend quotient method.

proc.c

found:

```
p->state = EMBRY0;
p->pid = nextpid++;
p->priority = 3;
```

When a process is created it is by default given a priority of 3.

```

int cps()
{
    struct proc *p;
    //Enables interrupts on this processor.
    sti();

    //Loop over process table looking for process with pid.
    acquire(&ptable.lock);
    cprintf("name \t pid \t state \t priority \n");
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state == SLEEPING)
            cprintf("%s \t %d \t SLEEPING \t %d \n ", p->name, p->pid, p->priority);
        else if(p->state == RUNNING)
            cprintf("%s \t %d \t RUNNING \t %d \n ", p->name, p->pid, p->priority);
        else if(p->state == RUNNABLE)
            cprintf("%s \t %d \t RUNNABLE \t %d \n ", p->name, p->pid, p->priority);
    }
    release(&ptable.lock);
    return 22;
}

```

Created a ps system call to list SLEEPING, RUNNING and RUNNABLE process.

```

int
nice(int pid, int priority)
{
    if(priority <1 || priority>5){
        cprintf("Error: Priority not within range(1-5)\n");
        exit();
    }
    struct proc *p;
    int found=0;
    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->pid == pid){
            found=1;
            break;
        }
    }
    if(!found){
        release(&ptable.lock);
        cprintf("No process with Pid=%d\n",pid);
        exit();
    }
    release(&ptable.lock);
    int old_priority = p->priority;
    p->priority=priority;

    remove_process_from_queue(p,old_priority);
    add_process_to_queue(p,priority);
    return (pid*MOD+old_priority);
}

```

The nice() function changes the priority of a process identified by pid to a new priority (1-5). It checks if the priority is valid, searches for the process in the process table, updates its priority, and adjusts the process's position in the scheduling queues accordingly. If the process is not found, it prints an error and exits.


```

struct node priority_q[PLEVELS][NPROC]; // A two dimensional array with 5 priority rows and 64 process columns
int marker[PLEVELS] = {0,0,0,0,0};      // 0 because initially no process inside the priority queue

int add_process_to_queue(struct proc *p, int pr)
{
    int i;
    for(i=0;i<marker[pr];i++){
        if(p->pid == priority_q[pr][i].p->pid){
            return -1;
        }
    }

    priority_q[pr][marker[pr]].p=p;
    marker[pr]++;
    return 1;
}

```

The `add_process_to_queue()` function adds a process `p` to a priority queue `priority_q` at the specified priority level `pr`. It first checks if the process is already in the queue, returning -1 if found. If not, it adds the process to the queue, increments the marker for that priority level, and returns 1 to indicate success.


```

int remove_process_from_queue(struct proc *p, int pr)
{
    int found = -1, i;
    for(i=0; i < marker[pr]; i++)
    {
        //if process is found
        if(priority_q[pr][i].p->pid == p->pid)
        {
            found = 1;
            break;
        }
    }
    if(found==-1)
    {
        return found;
    }
    int j;
    //shift process after removing
    for(j = i; j<marker[pr]-1; j++)
    {
        priority_q[pr][j] = priority_q[pr][j+1];
    }

    marker[pr]-=1;
    return found;
}

```

The `remove_process_from_queue()` function removes a process `p` from the priority queue at the specified priority level `pr`. It searches for the process in the queue, and if found, it shifts all subsequent processes one position to the left to maintain the queue order and decrements the marker for that priority level. It returns 1 if the process was found and removed, or -1 if the process was not found.

proc.h

```
C proc.h M X
C proc.h
1 // Per-CPU state
2 struct cpu {
3     uchar apicid;           // Local APIC ID
4     struct context *scheduler; // switch() here to enter scheduler
5     struct taskstate ts;     // Used by x86 to find stack for interrupt
6     struct segdesc gdt[NSEGS]; // x86 global descriptor table
7     volatile uint started;   // Has the CPU started?
8     int ncli;                // Depth of pushcli nesting.
9     int intena;              // Were interrupts enabled before pushcli?
10
11     // Cpu-local storage variables; see below
12     struct cpu *cpu;
13     struct proc *proc;       // The currently-running process.
14 };
15
16 struct node {
17     struct proc *p;
18     struct node *prev;
19     struct node *next;
20 };
21
22 extern struct cpu cpus[NCPU];
23 extern int ncpu;
24 extern int add_process_to_queue(struct proc *p, int priority);
25 extern int remove_remove_process_from_queueproc_from_q(struct proc *p, int priority);
26
```

Added definitions in proc.h

ps.c

```
C ps.c U X
C ps.c
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4 #include "fcntl.h"
5
6 int main(void){
7     cps();
8     exit();
9 }
```

Code for ps system call

syscall.c

```
Added lines ×
C syscall.c
83  extern int sys_exec(void);
84  extern int sys_exit(void);
85  extern int sys_fork(void);
86  extern int sys_fstat(void);
87  extern int sys_getpid(void);
88  extern int sys_kill(void);
89  extern int sys_link(void);
90  extern int sys_mkdir(void);
91  extern int sys_mknod(void);
92  extern int sys_open(void);
93  extern int sys_pipe(void);
94  extern int sys_read(void);
95  extern int sys_sbrk(void);
96  extern int sys_sleep(void);
97  extern int sys_unlink(void);
98  extern int sys_wait(void);
99  extern int sys_write(void);
100 extern int sys_uptime(void);
101 extern int sys_cps(void);
102 extern int sys_nice(void);
103
```

```
[SYS_close]    sys_close,
[SYS_cps]      sys_cps,
[SYS_nice]     sys_nice,
};
```

Added definitions in syscall.c

Syscall.h

```
C syscall.h M X
C syscall.h
1 // System call numbers
2 #define SYS_fork 1
3 #define SYS_exit 2
4 #define SYS_wait 3
5 #define SYS_pipe 4
6 #define SYS_read 5
7 #define SYS_kill 6
8 #define SYS_exec 7
9 #define SYS_fstat 8
10 #define SYS_chdir 9
11 #define SYS_dup 10
12 #define SYS_getpid 11
13 #define SYS_sbrk 12
14 #define SYS_sleep 13
15 #define SYS_uptime 14
16 #define SYS_open 15
17 #define SYS_write 16
18 #define SYS_mknod 17
19 #define SYS_unlink 18
20 #define SYS_link 19
21 #define SYS_mkdir 20
22 #define SYS_close 21
23 #define SYS_cps 22
24 #define SYS_nice 23
Changed lines
```

Gave system call numbers to the ps and nice calls

sysproc.c

C sysproc.c M X

C sysproc.c

```
61  {
78  }
79
80  // return how many clock tick inte
81  // since start.
82  int
83  sys_uptime(void)
84  {
85      uint xticks;
86
87      acquire(&tickslock);
88      xticks = ticks;
89      release(&tickslock);
90      return xticks;
91  }
92
93  int
94  sys_cps(void)
95  {
96      return cps();
97  }
98
99  int
100  sys_nice(void)
101  {
102      int pid, pr;
103      if(argint(0, &pid) < 0)
104          return -1;
105      if(argint(1, &pr) < 0)
106          return -1;
107
108      return nice(pid, pr);
109  }
```

Added process in sysproc

user.h

```
C user.h M X
C user.h
1  struct stat;
2  struct rtcddate;
3
4  // system calls
5  int fork(void);
6  int exit(void) __attribute__((noreturn));
7  int wait(void);
8  int pipe(int*);
9  int write(int, void*, int);
10 int read(int, void*, int);
11 int close(int);
12 int kill(int);
13 int exec(char*, char**);
14 int open(char*, int);
15 int mknod(char*, short, short);
16 int unlink(char*);
17 int fstat(int fd, struct stat*);
18 int link(char*, char*);
19 int mkdir(char*);
20 int chdir(char*);
21 int dup(int);
22 int getpid(void);
23 char* sbrk(int);
24 int sleep(int);
25 int uptime(void);
26 int cps(void);
27 int nice(int pid, int priority);
28
29
```

Added processes definitions in user.h

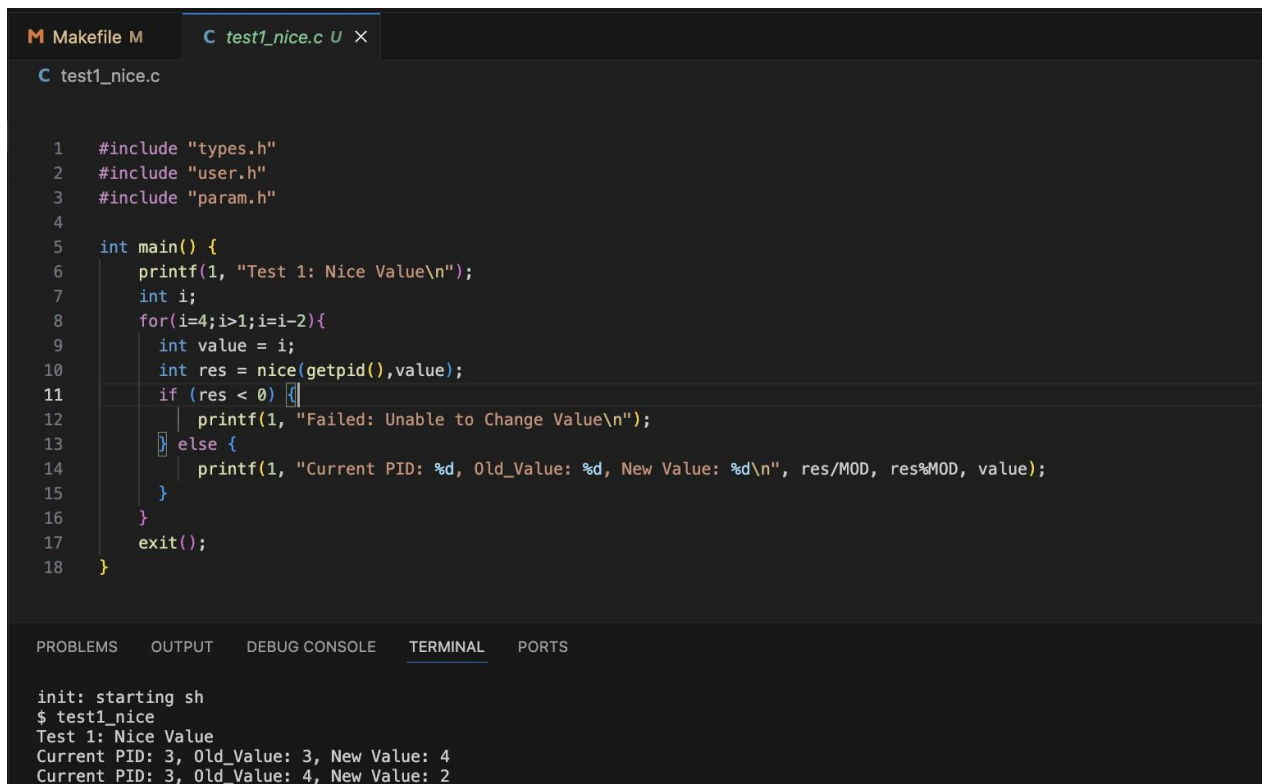
usys.S

```
ASM usys.S M X
ASM usys.S
4  #define SYSCALL(name) \
6  name: \
10
11  SYSCALL(fork)
12  SYSCALL(exit)
13  SYSCALL(wait)
14  SYSCALL(pipe)
15  SYSCALL(read)
16  SYSCALL(write)
17  SYSCALL(close)
18  SYSCALL(kill)
19  SYSCALL(exec)
20  SYSCALL(open)
21  SYSCALL(mknod)
22  SYSCALL(unlink)
23  SYSCALL(fstat)
24  SYSCALL(link)
25  SYSCALL(mkdir)
26  SYSCALL(chdir)
27  SYSCALL(dup)
28  SYSCALL(getpid)
29  SYSCALL(sbrk)
30  SYSCALL(sleep)
31  SYSCALL(uptime)
32  SYSCALL(cps)
33  SYSCALL(nice)
34  Added lines
```

Added syscalls.

TEST CASES

1. test1_nice.c



The image shows a code editor with two tabs: 'Makefile M' and 'C test1_nice.c U X'. The active tab 'test1_nice.c' displays the following C code:

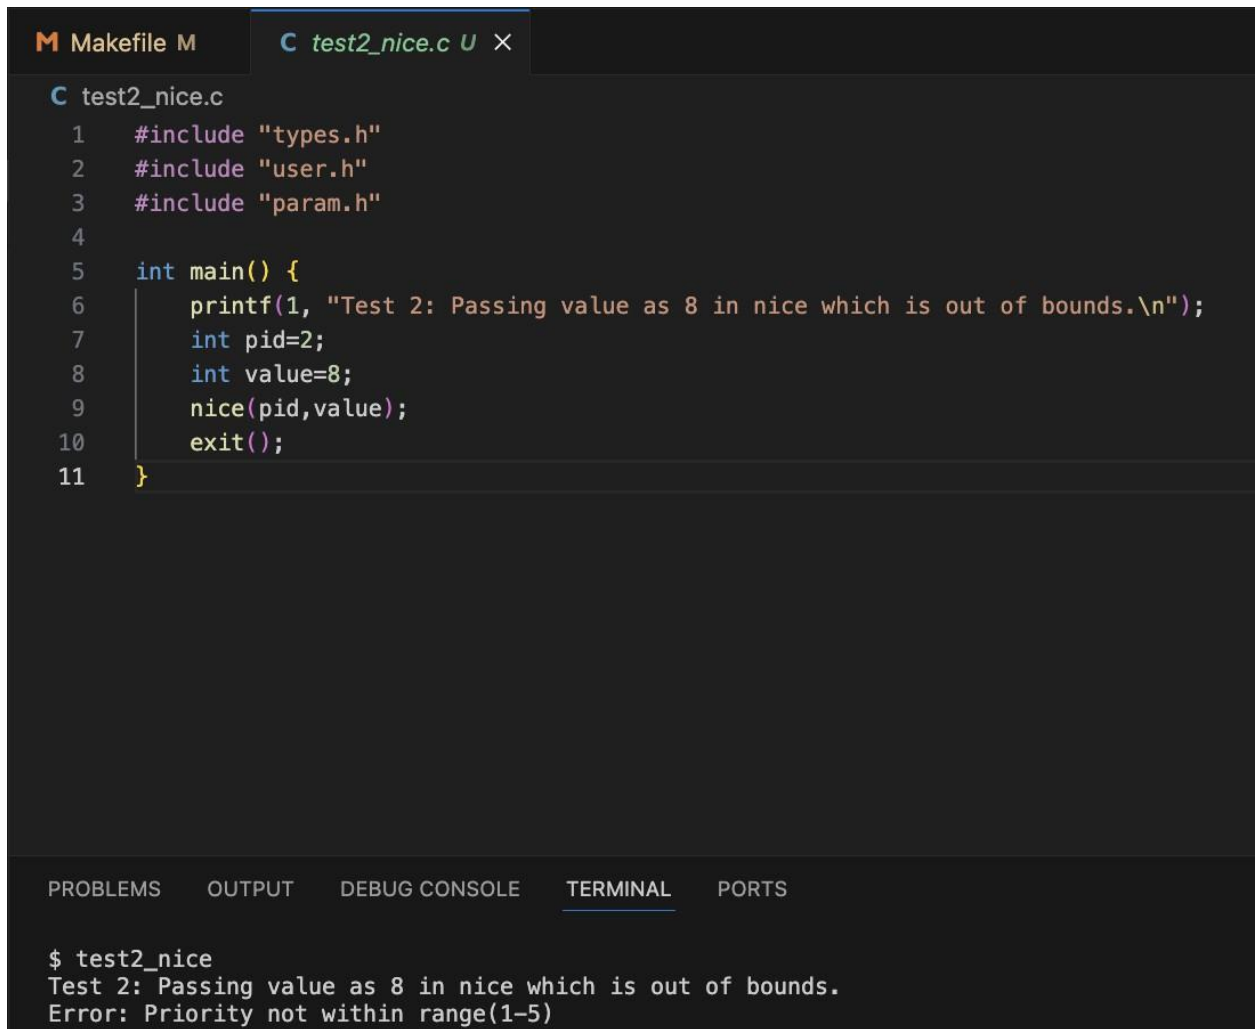
```
1  #include "types.h"
2  #include "user.h"
3  #include "param.h"
4
5  int main() {
6      printf(1, "Test 1: Nice Value\n");
7      int i;
8      for(i=4;i>1;i=i-2){
9          int value = i;
10         int res = nice(getpid(),value);
11         if (res < 0) {
12             printf(1, "Failed: Unable to Change Value\n");
13         } else {
14             printf(1, "Current PID: %d, Old_Value: %d, New Value: %d\n", res/MOD, res%MOD, value);
15         }
16     }
17     exit();
18 }
```

Below the code editor, there is a terminal window with the following output:

```
init: starting sh
$ test1_nice
Test 1: Nice Value
Current PID: 3, Old_Value: 3, New Value: 4
Current PID: 3, Old_Value: 4, New Value: 2
```

The program tests changing the priority of the current process twice using `nice()`, and it prints the process ID, old priority, and new priority value or an error if the change fails.

2. test2_nice.c



The image shows a code editor with two tabs: 'Makefile M' and 'C test2_nice.c U X'. The 'test2_nice.c' tab is active, displaying the following C code:

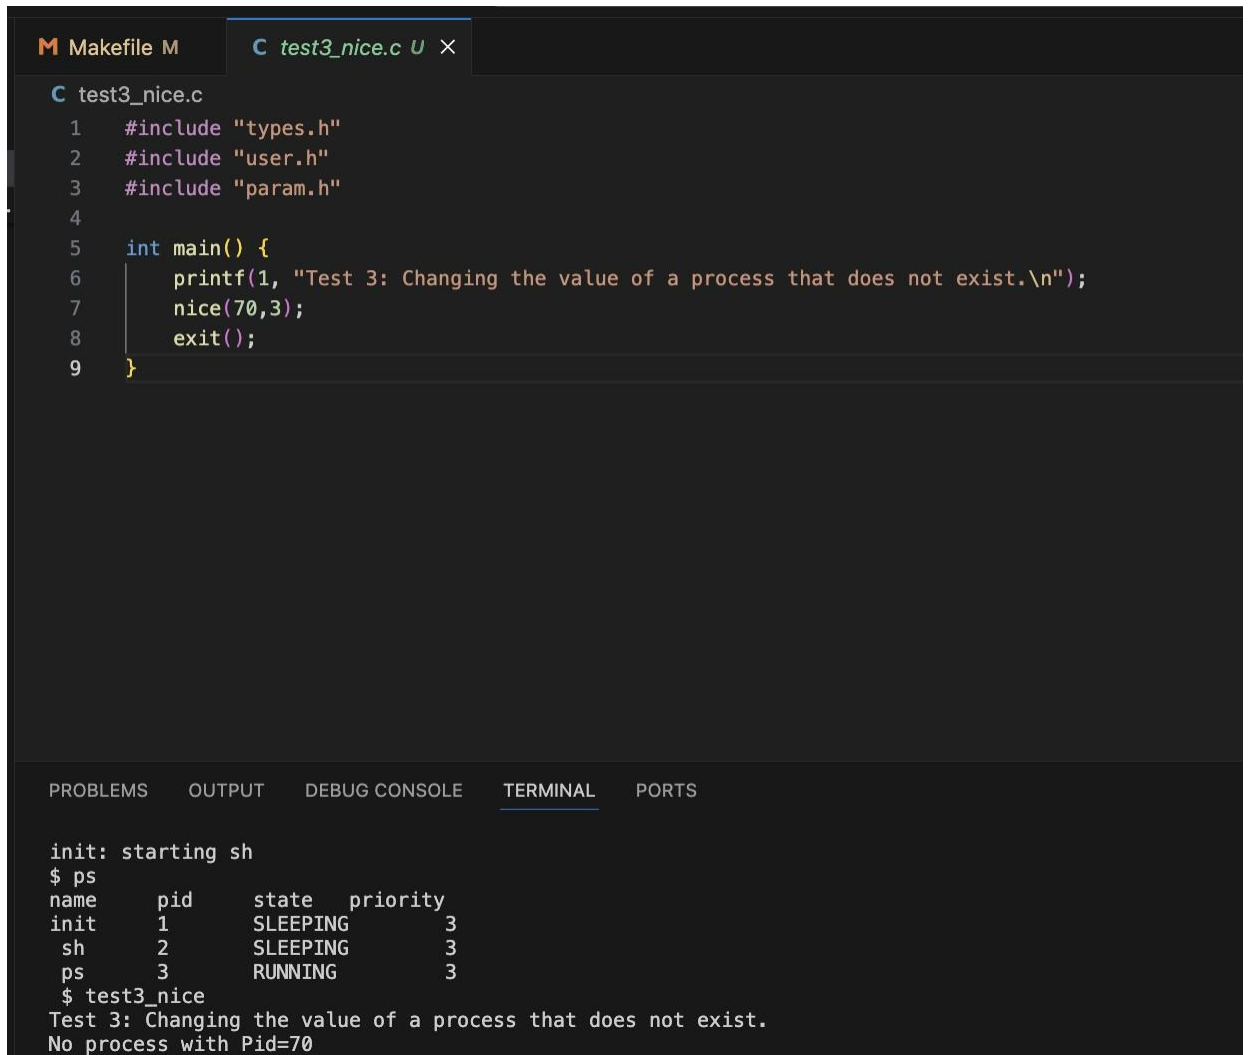
```
1  #include "types.h"
2  #include "user.h"
3  #include "param.h"
4
5  int main() {
6      printf(1, "Test 2: Passing value as 8 in nice which is out of bounds.\n");
7      int pid=2;
8      int value=8;
9      nice(pid,value);
10     exit();
11 }
```

Below the code editor, there is a terminal window with tabs: 'PROBLEMS', 'OUTPUT', 'DEBUG CONSOLE', 'TERMINAL', and 'PORTS'. The 'TERMINAL' tab is active, showing the execution of the program:

```
$ test2_nice
Test 2: Passing value as 8 in nice which is out of bounds.
Error: Priority not within range(1-5)
```

The program attempts to change the priority of a process with `pid = 2` to 8, which is out of bounds, likely triggering an error due to invalid priority.

3. test3_nice.c



The screenshot shows a code editor with two tabs: 'Makefile M' and 'C test3_nice.c U x'. The 'test3_nice.c' tab is active, displaying the following C code:

```
1 #include "types.h"
2 #include "user.h"
3 #include "param.h"
4
5 int main() {
6     printf(1, "Test 3: Changing the value of a process that does not exist.\n");
7     nice(70,3);
8     exit();
9 }
```

Below the code editor is a terminal window with tabs for 'PROBLEMS', 'OUTPUT', 'DEBUG CONSOLE', 'TERMINAL', and 'PORTS'. The 'TERMINAL' tab is active, showing the following output:

```
init: starting sh
$ ps
name    pid    state  priority
init     1    SLEEPING      3
sh       2    SLEEPING      3
ps       3    RUNNING       3
$ test3_nice
Test 3: Changing the value of a process that does not exist.
No process with Pid=70
```

The program tries to change the priority of a non-existent process with pid = 70 to 3, which will likely result in an error message indicating that the process was not found.

Task 2 Round Robin Priority Scheduler

MAKEFILE

M Makefile M X

C defs.h M

M Makefile

```
195 clean:
198     initcode initcode.out kernel xv6.img fs.img kernelmemfs mkfs \
199     .gdbinit \
200     $(UPROGS)
201
202 # make a printout
203 FILES = $(shell grep -v '^#' runoff.list)
204 PRINT = runoff.list runoff.spec README toc.hdr toc.ftr $(FILES)
205
206 xv6.pdf: $(PRINT)
207     ./runoff
208     ls -l xv6.pdf
209
210 print: xv6.pdf
211
212 # run in emulators
213
214 bochs : fs.img xv6.img
215     if [ ! -e .bochsrc ]; then ln -s dot-bochsrc .bochsrc; fi
216     bochs -q
217
218 # try to generate a unique GDB port
219 GDBPORT = $(shell expr `id -u` % 5000 + 25000)
220 # QEMU's gdb stub command line changed in 0.11
221 QEMUGDB = $(shell if $(QEMU) -help | grep -q '^-gdb'; \
222     then echo "-gdb tcp::$(GDBPORT)"; \
223     else echo "-s -p $(GDBPORT)"; fi)
224 ifndef CPUS
225 CPUS := 1
226 endif
227 QEMUOPTS = -drive file=fs.img,index=1,media=disk,format=raw -drive
228
```

Make CPU count to 1

```
M Makefile M X C defs.h M
M Makefile
57 # Try to infer the correct QEMU
58 ifndef QEMU
59 QEMU = $(shell if which qemu > /dev/null; \
60     then echo qemu; exit; \
61     elif which qemu-system-i386 > /dev/null; \
62     then echo qemu-system-i386; exit; \
63     else \
64     qemu=/Applications/Q.app/Contents/MacOS/i386-
65     if test -x $$qemu; then echo $$qemu; exit; fi
66     echo "***" 1>&2; \
67     echo "*** Error: Couldn't find a working QEMU
68     echo "*** Is the directory containing the qemu
69     echo "*** or have you tried setting the QEMU
70     echo "***" 1>&2; exit 1)
71 endif
72
73 CC = $(TOOLPREFIX)gcc
74 AS = $(TOOLPREFIX)gas
75 LD = $(TOOLPREFIX)ld
76 OBJCOPY = $(TOOLPREFIX)objcopy
77 OBJDUMP = $(TOOLPREFIX)objdump
78 CFLAGS = -fno-pic -static -fno-builtin -fno-strict-aliasing
79 #CFLAGS = -fno-pic -static -fno-builtin -fno-strict-aliasing
80 CFLAGS += $(shell $(CC) -fno-stack-protector -E ->
81 ASFLAGS = -m32 -gdwarf-2 -Wa,-divide
82 # FreeBSD ld wants ``elf_i386_fbsd''
83 LDFLAGS += -m $(shell $(LD) -V | grep elf_i386 2>
84 ifndef PRIORITY
85     PRIORITY := 1
86 endif
87 CFLAGS += -DCUSTOM_SCHEDULER=$(PRIORITY)
88
```

Add FLAG PRIORITY so that user has option to compile with normal RR and RR with priority.

proc.c

M Makefile M

C proc.c M X

C proc.c

```
93 found:
120
121     return p;
122 }
123
124 //PAGEBREAK: 32
125 // Set up first user process.
126 void
127 userinit(void)
128 {
129     struct proc *p;
130     extern char _binary_initcode_start[], _binary_initcode_size[];
131
132     p = allocproc();
133
134     initproc = p;
135     if((p->pgdir = setupkvm()) == 0)
136         panic("userinit: out of memory?");
137     inituvm(p->pgdir, _binary_initcode_start, (int)_binary_initcode_size);
138     p->sz = PGSIZE;
139     memset(p->tf, 0, sizeof(*p->tf));
140     p->tf->cs = (SEG_UCODE << 3) | DPL_USER;
141     p->tf->ds = (SEG_UDATA << 3) | DPL_USER;
142     p->tf->es = p->tf->ds;
143     p->tf->ss = p->tf->ds;
144     p->tf->eflags = FL_IF;
145     p->tf->esp = PGSIZE;
146     p->tf->eip = 0; // beginning of initcode.S
147
148     safestrcpy(p->name, "initcode", sizeof(p->name));
149     p->cwd = namei("/");
150
151     // this assignment to p->state lets other cores
152     // run this process. the acquire forces the above
153     // writes to be visible, and the lock is also needed
154     // because the assignment might not be atomic.
155     acquire(&ptable.lock);
156
157     p->state = RUNNABLE;
158
159     #ifndef CUSTOM_SCHEDULER
160     int res;
161     res=add_process_to_queue(p,p->priority-1);
162     if(res==-1){
163         cprintf("Unable to create process with pid=%d\n",p->pid);
164         release(&ptable.lock);
165         exit();
166     }
167     #endif
168     release(&ptable.lock);
169 }
```

The `userinit()` function initializes the first user process, setting up its memory, registers, and state as `RUNNABLE`, and, if using a custom scheduler, adds it to the appropriate priority queue. It handles synchronization with `ptable.lock` to ensure safe process state updates.

M Makefile M

C proc.c M X

C proc.c

```
192 // Sets up stack to return as if from system call.
193 // Caller must set state of returned proc to RUNNABLE.
194 int
195 fork(void)
196 {
197     int i, pid;
198     struct proc *np;
199
200     // Allocate process.
201     if((np = allocproc()) == 0){
202         return -1;
203     }
204
205     // Copy process state from p.
206     if((np->pgdir = copyvm(proc->pgdir, proc->sz)) == 0){
207         kfree(np->kstack);
208         np->kstack = 0;
209         np->state = UNUSED;
210         return -1;
211     }
212     np->sz = proc->sz;
213     np->parent = proc;
214     *np->tf = *proc->tf;
215
216     // Clear %eax so that fork returns 0 in the child.
217     np->tf->eax = 0;
218
219     for(i = 0; i < NOFILE; i++)
220         if(proc->ofile[i])
221             np->ofile[i] = filedup(proc->ofile[i]);
222     np->cwd = idup(proc->cwd);
223
224     safestrcpy(np->name, proc->name, sizeof(proc->name));
225
226     pid = np->pid;
227
228     acquire(&ptable.lock);
229
230     np->state = RUNNABLE;
231     #ifndef CUSTOM_SCHEDULER
232     int res;
233     res=add_process_to_queue(np,np->priority-1);
234     if(res==-1){
235         cprintf("Unable to create process with pid=%d\n",np->pid);
236         release(&ptable.lock);
237         exit();
238     }
239     #endif
240     release(&ptable.lock);
241     return pid;
242 }
```

This section checks if a custom scheduler is not being used. It then tries to add the new process `np` to the appropriate priority queue. If adding fails, it prints an error message and exits, ensuring proper handling of process creation failures.

C proc.c

```
288
289 // Wait for a child process to exit and return its pid.
290 // Return -1 if this process has no children.
291 int
292 wait(void)
293 {
294     struct proc *p;
295     int havekids, pid;
296
297     acquire(&ptable.lock);
298     for(;;){
299         // Scan through table looking for exited children.
300         havekids = 0;
301         for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
302             if(p->parent != proc)
303                 continue;
304             havekids = 1;
305             if(p->state == ZOMBIE){
306                 // Found one.
307                 pid = p->pid;
308                 kfree(p->kstack);
309                 p->kstack = 0;
310                 freevm(p->pgdir);
311                 p->pid = 0;
312                 p->parent = 0;
313                 p->name[0] = 0;
314                 p->killed = 0;
315                 p->state = UNUSED;
316                 #ifndef CUSTOM_SCHEDULER
317                 remove_process_from_queue(p, p->priority-1);
318                 #endif
319                 release(&ptable.lock);
320                 return pid;
321             }
322         }
323
324         // No point waiting if we don't have any children.
325         if(!havekids || proc->killed){
326             release(&ptable.lock);
327             return -1;
328         }
329
330         // Wait for children to exit. (See wakeup1 call in proc_exit.)
331         sleep(proc, &ptable.lock); //DOC: wait-sleep
332     }
333 }
```

In this part of the wait() function, if a custom scheduler is not in use, the code removes the process p from the priority queue when the process transitions to the UNUSED state after being freed. This ensures that terminated (ZOMBIE) processes are properly removed from scheduling queues before they are cleaned up.

```
struct cpu*
mycpu(void)
{
    int apicid, i;

    if(readeflags() & FL_IF)
        panic("mycpu called with interrupts enabled\n");

    apicid = lapic();
    // APIC IDs are not guaranteed to be contiguous. Maybe we should have
    // a reverse map, or reserve a register to store &cpus[i].
    for (i = 0; i < ncpu; ++i) {
        if (cpus[i].apicid == apicid)
            return &cpus[i];
    }
    panic("unknown apicid\n");
}

struct proc*
myproc(void) {
    struct cpu *c;
    struct proc *p;
    pushcli();
    c = mycpu();
    p = c->proc;
    popcli();
    return p;
}
```

mycpu(): This function returns a pointer to the cpu structure corresponding to the current CPU, identified by its APIC ID. It ensures interrupts are disabled before running and panics if the APIC ID is not recognized.

myproc(): This function returns a pointer to the currently running process on the current CPU. It temporarily disables interrupts to safely access the cpu and proc structures, then restores interrupt settings.

M Makefile M

C proc.c M X

C proc.c

```
371 // via switch back to the scheduler.
372 #ifndef CUSTOM_SCHEDULER
373 void
374 scheduler(void)
375 {
376     struct proc *p=0;
377     struct cpu *c = mycpu();
378     c->proc = 0;
379     for(;;){
380         // Enable interrupts on this processor.
381         sti();
382
383         // Loop over process table looking for process to run.
384         acquire(&ptable.lock);
385         int i;
386         for(i=0;i<PLEVELS;i++){
387             if(marker[i]>0){
388                 p=priority_q[i][0].p;
389                 remove_process_from_queue(p,i);
390                 break;
391             }
392         }
393
394         if(p!=0 && p->state==RUNNABLE){
395             c->proc = p;
396             switchvm(p);
397             p->state = RUNNING;
398             // Switch to chosen process. It is the process's job
399             // to release ptable.lock and then reacquire it
400             // before jumping back to us.
401             swtch(&(c->scheduler), p->context);
402             switchkvm();
403             // Process is done running for now.
404             // It should have changed its p->state before coming back.
405             c->proc = 0;
406         }
407         release(&ptable.lock);
408     }
409 }
410 #else
```

```

void
scheduler(void)
{
    struct proc *p;
    struct cpu *c = mycpu();
    c->proc = 0;

    for(;;){
        // Enable interrupts on this processor.
        sti();

        // Loop over process table looking for process to run.
        acquire(&ptable.lock);
        for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
            if(p->state != RUNNABLE)
                continue;

            // Switch to chosen process. It is the process's job
            // to release ptable.lock and then reacquire it
            // before jumping back to us.
            c->proc = p;
            switchvm(p);
            p->state = RUNNING;

            switch(&(c->scheduler), p->context);
            switchkvm();

            // Process is done running for now.
            // It should have changed its p->state before coming back.
            c->proc = 0;
        }
        release(&ptable.lock);
    }
}
#endif

```

The code defines two different versions of the scheduler() function, based on whether CUSTOM_SCHEDULER is defined:

1. **Without CUSTOM_SCHEDULER:**

- Implements a custom scheduling mechanism using priority queues.
- It iterates over the priority levels (PLEVELS), selects a process from the highest non-empty queue, removes it from the queue, and runs it if it is in the RUNNABLE state.

2. With CUSTOM_SCHEDULER:

- Implements a simpler round-robin scheduler.
- It iterates over all processes in the process table, selects a RUNNABLE process, and runs it.

In both versions, the scheduler enables interrupts, switches to the chosen process, and manages the process state transitions while holding the ptable.lock for safe access.

```
// Give up the CPU for one scheduling round.
void
yield(void)
{
    acquire(&ptable.lock); //DOC: yieldlock
    proc->state = RUNNABLE;
    #ifndef CUSTOM_SCHEDULER
        int res=add_process_to_queue(myproc(),myproc()->priority-1);
        if(res==-1){
            cprintf("Unable to create process with pid=%d\n",myproc()->pid);
        }
    #endif
    sched();
    release(&ptable.lock);
}
```

The yield() function makes the current process give up the CPU voluntarily, setting its state to RUNNABLE. If CUSTOM_SCHEDULER is not defined, it adds the process back to the appropriate priority queue. The sched() function is then called to perform a context switch, and finally, the ptable.lock is released.


```

//PAGEBREAK!
// Wake up all processes sleeping on chan.
// The ptable lock must be held.
static void
wakeup1(void *chan)
{
    struct proc *p;

    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->state == SLEEPING && p->chan == chan){
            p->state = RUNNABLE;
            #ifndef CUSTOM_SCHEDULER
            int res=add_process_to_queue(p,p->priority-1);
            if(res==-1){
                cprintf("Unable to create process with pid=%d\n",p->pid);
            }
            #endif
        }
    }
}

```

The wakeup1() function iterates over all processes in the process table and wakes up any process that is in the SLEEPING state and waiting on the specified chan. It changes the process state to RUNNABLE and, if CUSTOM_SCHEDULER is not defined, adds the process to the appropriate priority queue. If adding to the queue fails, it prints an error message.

```

int
kill(int pid)
{
    struct proc *p;

    acquire(&ptable.lock);
    for(p = ptable.proc; p < &ptable.proc[NPROC]; p++){
        if(p->pid == pid){
            p->killed = 1;
            // Wake process from sleep if necessary.
            if(p->state == SLEEPING)
                p->state = RUNNABLE;
            #ifndef CUSTOM_SCHEDULER
                int res;
                res=add_process_to_queue(p,p->priority-1);
                if(res==-1){
                    cprintf("Unable to create process with pid=%d\n",p->pid);
                    release(&ptable.lock);
                    exit();
                }
            #endif
            release(&ptable.lock);
            return 0;
        }
    }
    release(&ptable.lock);
    return -1;
}

```

The kill() function sets the killed flag for a process with the specified pid and wakes it up if it is SLEEPING by setting its state to RUNNABLE. If CUSTOM_SCHEDULER is not defined, it adds the process to the priority queue. If adding fails, it prints an error and exits. The function returns 0 on success or -1 if no process with the given pid is found.

HOW TO RUN SCHEDULER WITH PRIORITY RR

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```
$ make PRIORITY=1
dd if=/dev/zero of=xv6.img count=10000
10000+0 records in
10000+0 records out
5120000 bytes (5.1 MB) copied, 0.0313446 s, 163 MB/s
dd if=bootblock of=xv6.img conv=notrunc
1+0 records in
1+0 records out
512 bytes (512 B) copied, 0.0015145 s, 338 kB/s
dd if=kernel of=xv6.img seek=1 conv=notrunc
360+1 records in
360+1 records out
184364 bytes (184 kB) copied, 0.00220762 s, 83.5 MB/s
$ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 1 -m 512
xv6...
cpu0: starting
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ █
```

TEST CASES

C test1.c U X

C test1.c

```
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4
5  void prime() {
6      int num = 2; // Start from 2, since 1 is not a prime number
7      while (num < 7) {
8          int prime = 1;
9          int i;
10         for (i = 2; i * i <= num; i++) { // Corrected condition to i * i <= num
11             if (num % i == 0) {
12                 prime = 0;
13                 break;
14             }
15         }
16         if (prime) {
17             printf(1, "Pid: %d, Prime: %d\n", getpid(), num);
18         }
19         num++;
20     }
21 }
22
23
24 int main() {
25     printf(1, "----- Test Case 1 with Custom Scheduler ----- \n");
26     printf(1, "Lowering the Parent's Priority Below the Child's Priority \n");
27     int pid = fork();
28     sleep(200); // to process changing the pid
29     if (pid < 0) {
30         printf(2, "Fork fail in pid: %d\n", getpid());
31         exit();
32     }
33     if (pid > 0) {
34         printf(1, "Change Parents priority to 1 with pid=%d\n", getpid());
35         nice(getpid(), 1);
36     }
37     if (pid == 0) {
38         printf(1, "Child process with pid=%d\n", getpid());
39         // nice(getpid(), 5);
40         prime();
41         exit();
42     } else {
43         printf(1, "Parent process with pid=%d\n", getpid());
44         prime();
45         wait();
46         exit();
47     }
48     exit();
49 }
50
```

1.

```

31     exit();
32 }
33 if (pid > 0) {
34     printf(1, "Change Parents priority to 1 with pid=%d\n", getpid());
35     nice(getpid(), 1);
36 }

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```

360+1 records out
184364 bytes (184 kB) copied, 0.00220762 s, 83.5 MB/s
$ make qemu-nox
qemu-system-i386 -nographic -drive file=fs.img,index=1,media=disk,format=raw -drive file=xv6.img,index=0,media=disk,format=raw -smp 1 -m 512
xv6...
cpu0: starting
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test1
----- Test Case 1 with Custom Scheduler -----
Lowering the Parent's Priority Below the Child's Priority
Change Parents priority to 1 with pid=3
Parent process with pid=3
Pid: 3, Prime: 2
Pid: 3, Prime: 3
Pid: 3, Prime: 5
Child process with pid=4
Pid: 4, Prime: 2
Pid: 4, Prime: 3
Pid: 4, Prime: 5
$

```

The code tests a custom scheduler by creating a parent and child process. By default, the child process has a higher priority, but the parent's priority is lowered and executes first.

C test2.c U X

C test2.c

```
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4
5  // Function to calculate and print prime numbers
6  void find_primes(int priority, int t) {
7      int pid = getpid();
8      // Set the process priority using the nice system call
9      nice(pid, priority);
10
11     int start_time = uptime();
12     int i = 1;
13
14     printf(1, "Started Current PID: %d, Priority: %d\n", pid, priority);
15
16     while (1) {
17         int j;
18
19         // Check if the number is a prime
20         for (j = 2; j * j <= i; j++) {
21             if (i % j == 0) {
22                 break;
23             }
24         }
25
26         int current_time = uptime();
27         if (current_time - start_time >= t) {
28             int end_time = uptime();
29             int execution_time = end_time - start_time;
30             printf(1, "Finished Current PID: %d, Priority: %d, Execution Time: %d\n", pid, priority, execution_time);
31             exit();
32         }
33         i++;
34     }
35 }
36
37 int main() {
38     int max_time = 1000;
39
40     // Creating 3 children with same priority
41     int p1 = fork();
42     if (p1 == 0) {
43         find_primes(2, max_time);
44         exit();
45     }
46
47     int p2 = fork();
48     if (p2 == 0) {
49         find_primes(2, max_time);
50         exit();
51     }
52
53     int p3 = fork();
54     if (p3 == 0) {
55         find_primes(2, max_time);
56         exit();
57     }
58     while(wait() != -1);
59     exit();
60 }
61
62
```

2.

```

22         break;
23     }
24 }
25
26 int current_time = uptime();
27 if (current_time - start_time >= t) {
28     int end_time = uptime();
29     int execution_time = end_time - start_time;
30     printf(1, "Finished Current PID: %d, Priority: %d, Execution Time: %d\n", pid, priority, execution_time);
31     exit();
32 }
33 i++;
34 }
35 }

```

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS

```

$ test2
Started Current PID: 25, Priority: 2
Started Current PID: 26, Priority: 2
Started Current PID: 27, Priority: 2
Finished Current PID: 25, Priority: 2, Execution Time: 1000
Finished Current PID: 26, Priority: 2, Execution Time: 1001
Finished Current PID: 27, Priority: 2, Execution Time: 1000
$

```

The code creates three child processes, each with the same priority level (2). Each child process calculates prime numbers for a specified duration (max_time = 1000 ticks) and then exits. Since all three processes have identical priority, they share the CPU time equally, resulting in similar execution times for each process. This demonstrates fair scheduling when processes have the same priority.


```

C test3.c U X
C test3.c
1  #include "types.h"
2  #include "stat.h"
3  #include "user.h"
4
5  // Function to calculate and print prime numbers up to 1000
6  void calculate_primes(int priority, int t) {
7      int pid = getpid();
8      // Set the process priority using the nice system call
9      nice(pid,priority);
10     int i,j;
11     int start_time = uptime();
12     int prime_count = 0;
13
14     printf(1, "Started Current PID: %d, Priority: %d\n", pid, priority);
15
16     // Loop to find all prime numbers up to 1000
17     for (i = 2; i <= 1000000; i++) {
18         int is_prime = 1;
19         for (j = 2; j * j <= i; j++) {
20             if (i % j == 0) {
21                 is_prime = 0;
22                 break;
23             }
24         }
25         if (is_prime) {
26             prime_count++;
27         }
28     }
29
30     int end_time = uptime();
31     int execution_time = end_time - start_time;
32
33     printf(1, "Finished Current PID: %d, Priority: %d, Execution Time: %d ticks, Total Primes Found: %d\n", pid, priority, execution_time, prime_count);
34
35     exit();
36 }
37
38 // Main function to run multiple processes
39 int main(int argc, char *argv[]) {
40     int priorities[] = {2, 2, 4, 4}; // Example priorities for each process
41     int process_count = sizeof(priorities) / sizeof(priorities[0]);
42     int t = 100; // Time threshold in ticks
43     int i;
44     for (i = 0; i < process_count; i++) {
45         if (fork() == 0) {
46             // Child process: calculate primes
47             calculate_primes(priorities[i], t);
48         }
49     }
50
51     // Wait for all child processes to finish
52     for (i = 0; i < process_count; i++) {
53         wait();
54     }
55 }

```

3.

```
C test3.c U X C test2.c U C test1.c U M Makefile M
C test3.c
1 #include "types.h"
2 #include "stat.h"
3 #include "user.h"
4
5 // Function to calculate and print prime numbers up to 1000
6 void calculate_primes(int priority, int t) {
7     int pid = getpid();
8     // Set the process priority using the nice system call
9     nice(pid,priority);
10    int i,j;
11    int start_time = uptime();
12    int prime_count = 0;
13
14    printf(1, "Started Current PID: %d, Priority: %d\n", pid, priority);
15
16    // Loop to find all prime numbers up to 1000
17    for (i = 2; i <= 1000000; i++) {
18        int is_prime = 1;
19        for (j = 2; j * j <= i; j++) {
20            if (i % j == 0) {
21                is_prime = 0;
22                break;
23            }
24        }
25        if (is_prime) {
26            prime_count++;
27        }
28    }
29
30    int end_time = uptime();
31    int execution_time = end_time - start_time;
32
33    printf(1, "Finished Current PID: %d, Priority: %d, Execution Time: %d ticks, Total Primes Found: %d\n", pid, priority, execution_time, prime_count);
34}

PROBLEMS OUTPUT DEBUG CONSOLE TERMINAL PORTS
cpu0: starting
sb: size 1000 nblocks 941 ninodes 200 nlog 30 logstart 2 inodestart 32 bmap start 58
init: starting sh
$ test3
Started Current PID: 4, Priority: 2
Finished Current PID: 4, Priority: 2, Execution Time: 261 ticks, Total Primes Found: 78498
Started Current PID: 5, Priority: 2
Finished Current PID: 5, Priority: 2, Execution Time: 233 ticks, Total Primes Found: 78498
Started Current PID: 6, Priority: 4
Started Current PID: 7, Priority: 4
Finished Current PID: 6, Priority: 4, Execution Time: 463 ticks, Total Primes Found: 78498
Finished Current PID: 7, Priority: 4, Execution Time: 469 ticks, Total Primes Found: 78498
$
```

The code creates multiple child processes, each with different priorities: two processes with priority 2 (higher priority) and two with priority 4 (lower priority). Because processes with lower numerical priority values (like 2) are given higher scheduling priority, they are executed first. As a result, these higher-priority processes complete their prime number calculations faster, leading to shorter execution times compared to the lower-priority processes.